

Time-Series Analysis and Forecasting with Corn and Wheat Prices

Part 2 - Generate Seasonal ARIMA Forecasting Model for Corn Price Time-Series Data

A. Import Required Libraries

Import numpy and pandas libraries, and set %matplotlib inline.

```
In [1]: import numpy as np
import pandas as pd
%matplotlib inline
```

Import seasonal decomposition, autocorrelation (ACF) and partial autocorrelation function (PACF) graphs, differencing, augmented Dickey-Fuller test, and seasonal ARIMA with exogenous regressors (SARIMAX) tools from statsmodels library.

```
In [2]: from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.statespace.tools import diff
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

Import Pyramid Arima (pmdarima) tool to determine seasonal ARIMA (SARIMA) orders and select best SARIMA model.

```
In [3]: from pmdarima.arima import auto_arima
```

Load specific evaluation tools.

```
In [4]: from sklearn.metrics import mean_squared_error
from statsmodels.tools.eval_measures import rmse
```

Ignore harmless warnings.

```
In [5]: import warnings
warnings.filterwarnings('ignore')
```

B. Import, Inspect, and Visualize Joined Corn and Wheat Price Time-Series Data Set

Import corn and wheat price data pickle file into Pandas dataframe called corn_wheat2.

```
In [6]: corn_wheat2 = pd.read_pickle('C://Users/kyrma/Python_for_TS_DA/Capstone_Project/')
```

Check number of rows and columns in corn_wheat2 dataframe.

```
In [7]: corn_wheat2.shape
```

```
Out[7]: (450, 2)
```

View structure of corn_wheat2 dataframe.

```
In [8]: corn_wheat2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 450 entries, 1980-01-01 to 2017-06-01
Data columns (total 2 columns):
Corn_Price      450 non-null float64
Wheat_Price     450 non-null float64
dtypes: float64(2)
memory usage: 10.5 KB
```

View first five rows of corn_wheat2 dataframe.

```
In [9]: corn_wheat2.head()
```

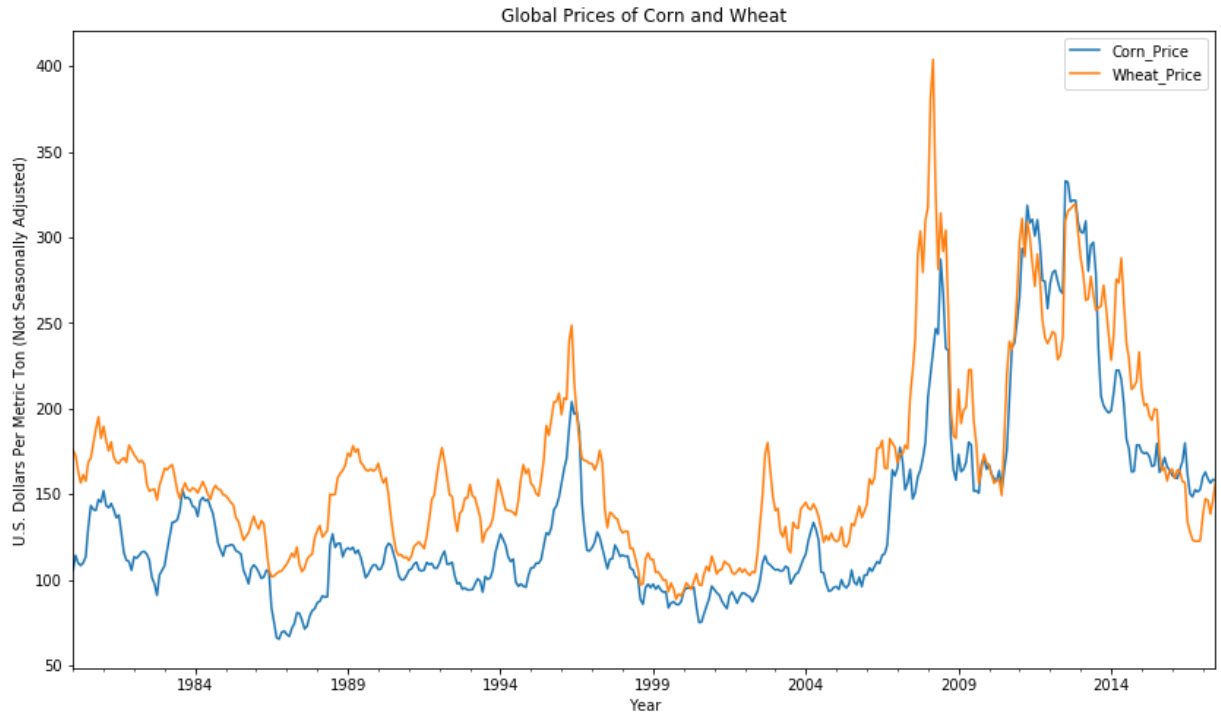
```
Out[9]:
```

	Corn_Price	Wheat_Price
Date		
1980-01-01	105.506813	175.634750
1980-02-01	114.167831	172.695236
1980-03-01	109.837318	163.509323
1980-04-01	108.262604	156.528030
1980-05-01	109.837318	161.304703

Plot out the corn and wheat price data with a reasonable figure size.

```
In [10]: title1 = 'Global Prices of Corn and Wheat'
ylabel1 = 'U.S. Dollars Per Metric Ton (Not Seasonally Adjusted)'
xlabel1 = 'Year'

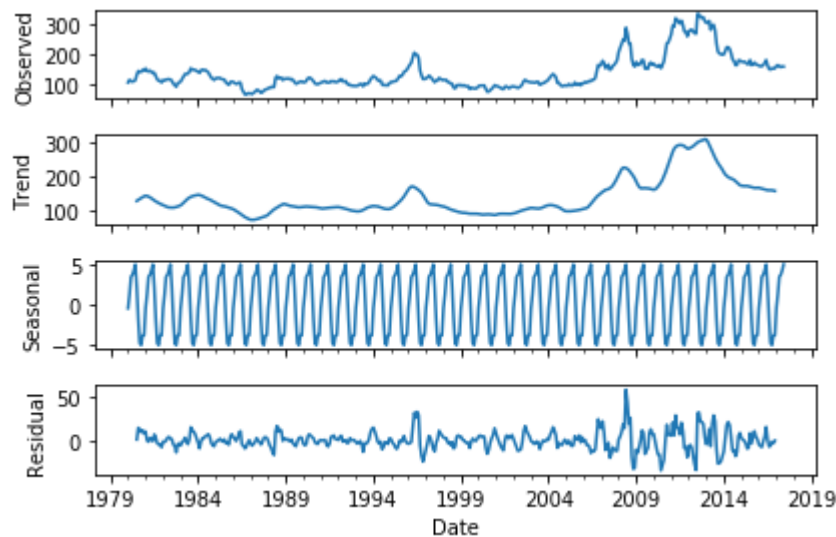
ax1 = corn_wheat2.plot(figsize=(14,8), title=title1)
ax1.autoscale(axis='x', tight=True)
ax1.set(xlabel=xlabel1, ylabel=ylabel1);
```



C. Build SARIMA Model and Forecast One Year Out for Corn Price Time-Series Data

Run an additive error, trend, and seasonality (ETS) decomposition on the global corn price data. The decomposition is additive because the corn price time-series data does not show any exponential growth.

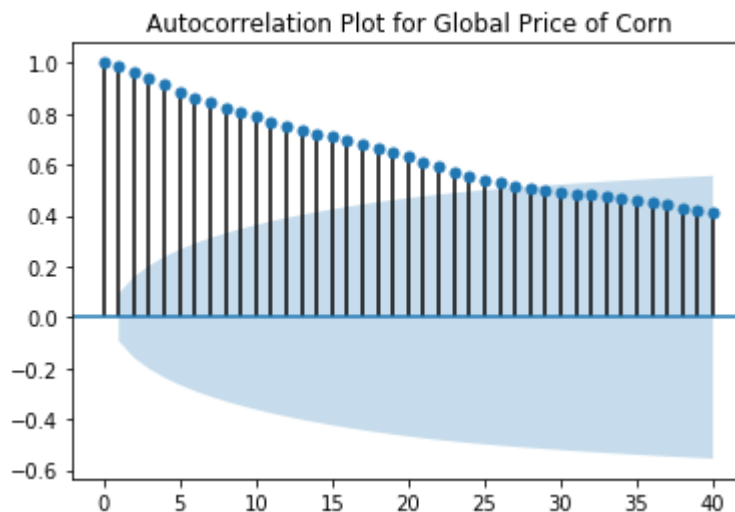
```
In [11]: corn_result = seasonal_decompose(corn_wheat2['Corn_Price'], model='add')
corn_result.plot();
```



- The additive ETS decomposition on the global corn price data reveals that the trend component runs a similar course to the observed time series. Moreover, the time series has a seasonal component and a large residual component. Even though the seasonal component is small in scale compared to the overall values, it exhibits a definite annual seasonality. Monthly global corn prices fluctuate from planting and harvest seasons in corn-growing countries, planted acreage, weather, and growing conditions.

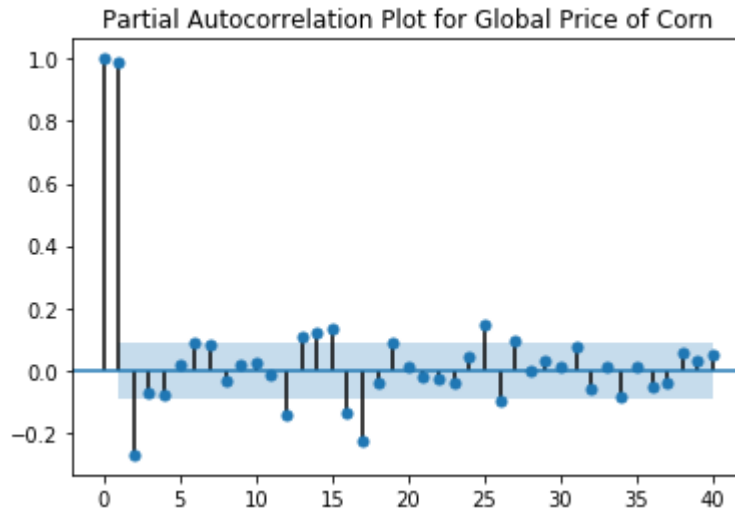
Plot ACF plot to determine if the global corn price data is stationary or not.

```
In [12]: title2 = 'Autocorrelation Plot for Global Price of Corn'
lags1 = 40
plot_acf(corn_wheat2['Corn_Price'], title=title2, lags=lags1);
```



Plot PACF plot to determine if the global corn price data is stationary or not.

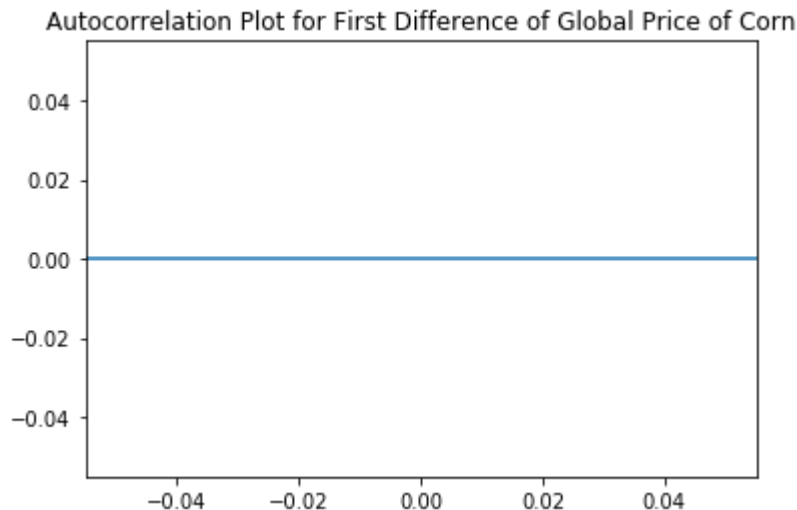
```
In [13]: title3 = 'Partial Autocorrelation Plot for Global Price of Corn'
plot_pacf(corn_wheat2['Corn_Price'], title=title3, lags=lags1);
```



- Based on the above ACF plot, the global corn price data is not stationary. Therefore, the data needs to be differenced until it is stationary.

Plot ACF plot of first difference for global corn price data to determine if it is stationary or not.

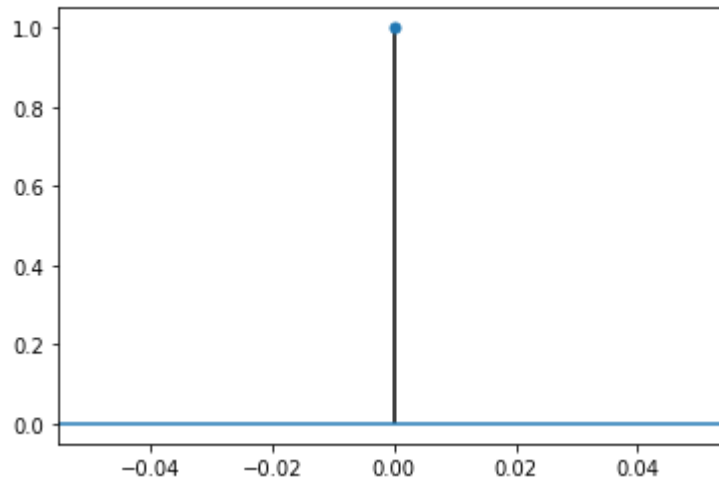
```
In [14]: corn_wheat2['Corn_Price_D1'] = diff(corn_wheat2['Corn_Price'], k_diff=1)
title4 = 'Autocorrelation Plot for First Difference of Global Price of Corn'
plot_acf(corn_wheat2['Corn_Price_D1'], title=title4, lags=lags1);
```



Plot PACF plot of first difference for global corn price data to determine if it is stationary or not.

```
In [15]: title5 = 'Partial Autocorrelation Plot for First Difference of Global Price of Corn'
plot_pacf(corn_wheat2['Corn_Price_D1'], title=title5, lags=lags1);
```

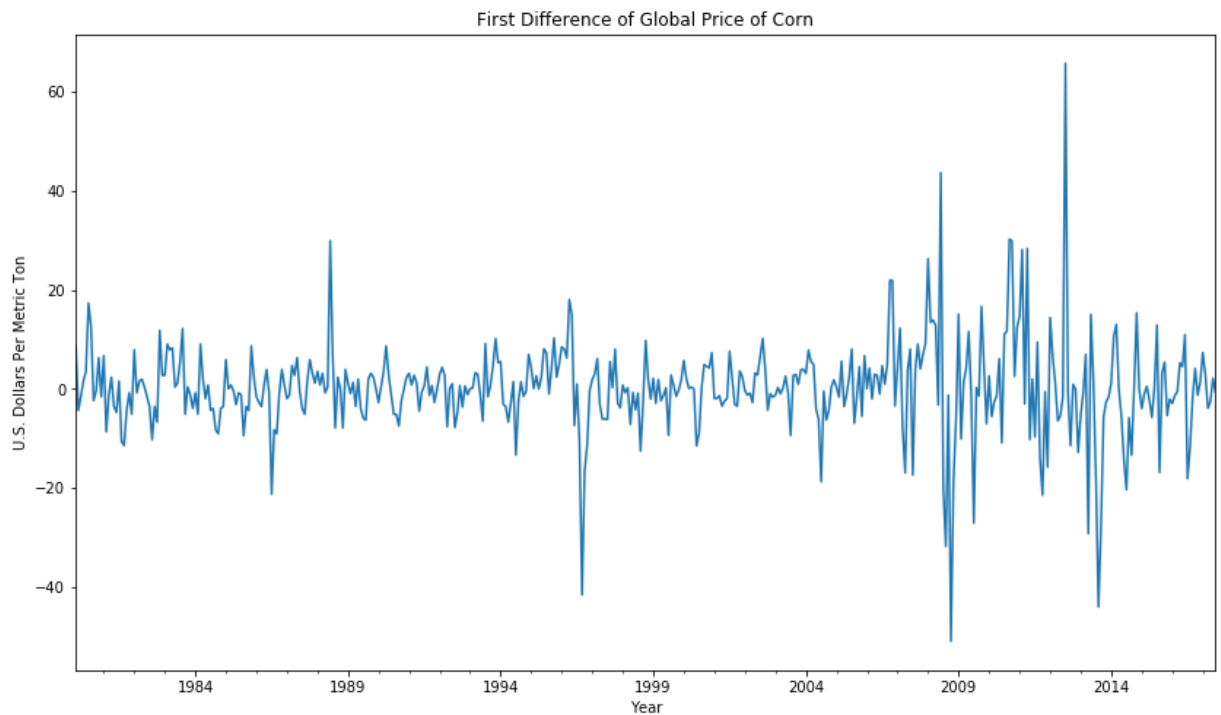
Partial Autocorrelation Plot for First Difference of Global Price of Corn



Plot out first difference of global corn price data with a reasonable figure size.

```
In [16]: title6 = 'First Difference of Global Price of Corn'
ylabel2 = 'U.S. Dollars Per Metric Ton'
xlabel2 = 'Year'

ax2 = corn_wheat2['Corn_Price_D1'].plot(figsize=(14,8), title=title6)
ax2.autoscale(axis='x', tight=True)
ax2.set(xlabel=xlabel2, ylabel=ylabel2);
```



- Based on the above PACF and first difference of global corn price data, the first difference is stationary.

Activate following function for running augmented Dickey-Fuller test.

```
In [17]: def adf_test(series, title=''):
        """
        Pass in a time series and an optional title, returns an ADF report
        """
        print(f'Augmented Dickey-Fuller Test: {title}')
        result = adfuller(series.dropna(), autolag='AIC') # .dropna() handles differ

        labels = ['ADF test statistic', 'p-value', '# lags used', '# observations']
        out = pd.Series(result[0:4], index=labels)

        for key, val in result[4].items():
            out[f'critical value ({key})']=val

        print(out.to_string()) # .to_string() removes the line "dtype: float64"

        if result[1] <= 0.05:
            print("Strong evidence against the null hypothesis")
            print("Reject the null hypothesis")
            print("Data has no unit root and is stationary")
        else:
            print("Weak evidence against the null hypothesis")
            print("Fail to reject the null hypothesis")
            print("Data has a unit root and is non-stationary")
```

Subject global corn price data to augmented Dickey-Fuller test.

```
In [18]: adf_test(corn_wheat2['Corn_Price'], title='Global Price of Corn')
```

```
Augmented Dickey-Fuller Test: Global Price of Corn
ADF test statistic      -2.150161
p-value                 0.224828
# lags used             16.000000
# observations          433.000000
critical value (1%)     -3.445543
critical value (5%)     -2.868238
critical value (10%)    -2.570338
Weak evidence against the null hypothesis
Fail to reject the null hypothesis
Data has a unit root and is non-stationary
```

Subject first difference of global corn price data to augmented Dickey-Fuller test.

```
In [19]: adf_test(corn_wheat2['Corn_Price_D1'], title='First Difference of Global Price of Corn')
```

```
Augmented Dickey-Fuller Test: First Difference of Global Price of Corn
ADF test statistic      -4.976504
p-value                0.000025
# lags used            15.000000
# observations         433.000000
critical value (1%)    -3.445543
critical value (5%)    -2.868238
critical value (10%)   -2.570338
Strong evidence against the null hypothesis
Reject the null hypothesis
Data has no unit root and is stationary
```

Run `pmdarima.auto_arima`'s stepwise function to obtain recommended orders of (p,d,q) and (P,D,Q)m for SARIMA model based off smallest Akaike information criterion (AIC). P, D, and Q represent the seasonal regression, differencing, and moving average coefficients, respectively, while m represents the number of rows in each seasonal cycle. The SARIMA model is built from the entire sample of global corn price data.


```
In [20]: corn_stepwise_fit = auto_arima(corn_wheat2['Corn_Price'], start_p=0, start_q=0,
                                         max_p=6, max_q=3, m=12,
                                         seasonal=True,
                                         d=None, trace=True,
                                         error_action='ignore', # We don't want to know if
                                         suppress_warnings=True, # We don't want convergen
                                         stepwise=True) # Set to stepwise.

corn_stepwise_fit.summary()
```

```
Fit ARIMA: order=(0, 1, 0) seasonal_order=(1, 0, 1, 12); AIC=3295.496, BIC=331
1.924, Fit time=0.230 seconds
Fit ARIMA: order=(0, 1, 0) seasonal_order=(0, 0, 0, 12); AIC=3295.484, BIC=330
3.698, Fit time=0.013 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(1, 0, 0, 12); AIC=3273.262, BIC=328
9.691, Fit time=0.254 seconds
Fit ARIMA: order=(0, 1, 1) seasonal_order=(0, 0, 1, 12); AIC=3275.408, BIC=329
1.836, Fit time=0.287 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(0, 0, 0, 12); AIC=3272.207, BIC=328
4.528, Fit time=0.039 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(0, 0, 1, 12); AIC=3272.914, BIC=328
9.342, Fit time=0.266 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(1, 0, 1, 12); AIC=3272.153, BIC=329
2.688, Fit time=0.340 seconds
Fit ARIMA: order=(2, 1, 0) seasonal_order=(1, 0, 1, 12); AIC=3273.325, BIC=329
7.967, Fit time=0.663 seconds
Fit ARIMA: order=(1, 1, 1) seasonal_order=(1, 0, 1, 12); AIC=3272.944, BIC=329
7.586, Fit time=0.650 seconds
Fit ARIMA: order=(2, 1, 1) seasonal_order=(1, 0, 1, 12); AIC=nan, BIC=nan, Fit
time=nan seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 0, 1, 12); AIC=3261.451, BIC=328
6.093, Fit time=1.477 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 0, 0, 12); AIC=3266.170, BIC=328
6.705, Fit time=0.962 seconds
Fit ARIMA: order=(1, 1, 0) seasonal_order=(2, 0, 2, 12); AIC=3262.983, BIC=329
1.732, Fit time=2.511 seconds
Fit ARIMA: order=(0, 1, 0) seasonal_order=(2, 0, 1, 12); AIC=3285.645, BIC=330
6.180, Fit time=1.106 seconds
Fit ARIMA: order=(2, 1, 0) seasonal_order=(2, 0, 1, 12); AIC=3262.565, BIC=329
1.314, Fit time=2.317 seconds
Fit ARIMA: order=(1, 1, 1) seasonal_order=(2, 0, 1, 12); AIC=3262.083, BIC=329
0.832, Fit time=1.669 seconds
Fit ARIMA: order=(2, 1, 1) seasonal_order=(2, 0, 1, 12); AIC=nan, BIC=nan, Fit
time=nan seconds
Total fit time: 12.809 seconds
```

Out[20]: Statespace Model Results

Dep. Variable:	y	No. Observations:	450
Model:	SARIMAX(1, 1, 0)x(2, 0, 1, 12)	Log Likelihood	-1624.725
Date:	Wed, 22 May 2019	AIC	3261.451
Time:	17:55:33	BIC	3286.093
Sample:	0	HQIC	3271.164

- 450

Covariance Type:

opg

	coef	std err	z	P> z	[0.025	0.975]
intercept	0.2838	0.741	0.383	0.702	-1.169	1.736
ar.L1	0.2397	0.030	8.074	0.000	0.182	0.298
ar.S.L12	-0.7519	0.122	-6.146	0.000	-0.992	-0.512
ar.S.L24	-0.1633	0.040	-4.035	0.000	-0.243	-0.084
ma.S.L12	0.7241	0.117	6.208	0.000	0.495	0.953
sigma2	81.1386	2.124	38.206	0.000	76.976	85.301

Ljung-Box (Q): 48.19 **Jarque-Bera (JB):** 2748.78**Prob(Q):** 0.18 **Prob(JB):** 0.00**Heteroskedasticity (H):** 5.76 **Skew:** 0.45**Prob(H) (two-sided):** 0.00 **Kurtosis:** 15.09

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Split global corn and wheat price data into training and test sets. The test set is 12 rows long since I want to make a year-long forecast.

```
In [21]: train = corn_wheat2.iloc[:len(corn_wheat2) - 12]
test = corn_wheat2.iloc[len(corn_wheat2) - 12:]
```

Fit an SARIMA(1,1,0)(2,0,1,12) model to global corn price training set.

```
In [22]: corn_train_model = SARIMAX(train['Corn_Price'], order=(1,1,0), seasonal_order=(2,
corn_results = corn_train_model.fit()
corn_results.summary()
```

Out[22]: Statespace Model Results

Dep. Variable:	Corn_Price	No. Observations:	438
Model:	SARIMAX(1, 1, 0)x(2, 0, 1, 12)	Log Likelihood	-1584.957
Date:	Wed, 22 May 2019	AIC	3179.915
Time:	17:55:34	BIC	3200.314
Sample:	01-01-1980	HQIC	3187.964
	- 06-01-2016		
Covariance Type:	opg		

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.2392	0.031	7.713	0.000	0.178	0.300
ar.S.L12	-0.6831	0.181	-3.777	0.000	-1.037	-0.329
ar.S.L24	-0.1614	0.041	-3.945	0.000	-0.242	-0.081
ma.S.L12	0.6529	0.173	3.764	0.000	0.313	0.993
sigma2	82.5526	2.184	37.805	0.000	78.273	86.832

Ljung-Box (Q):	47.84	Jarque-Bera (JB):	2623.98
Prob(Q):	0.18	Prob(JB):	0.00
Heteroskedasticity (H):	5.96	Skew:	0.46
Prob(H) (two-sided):	0.00	Kurtosis:	14.97

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

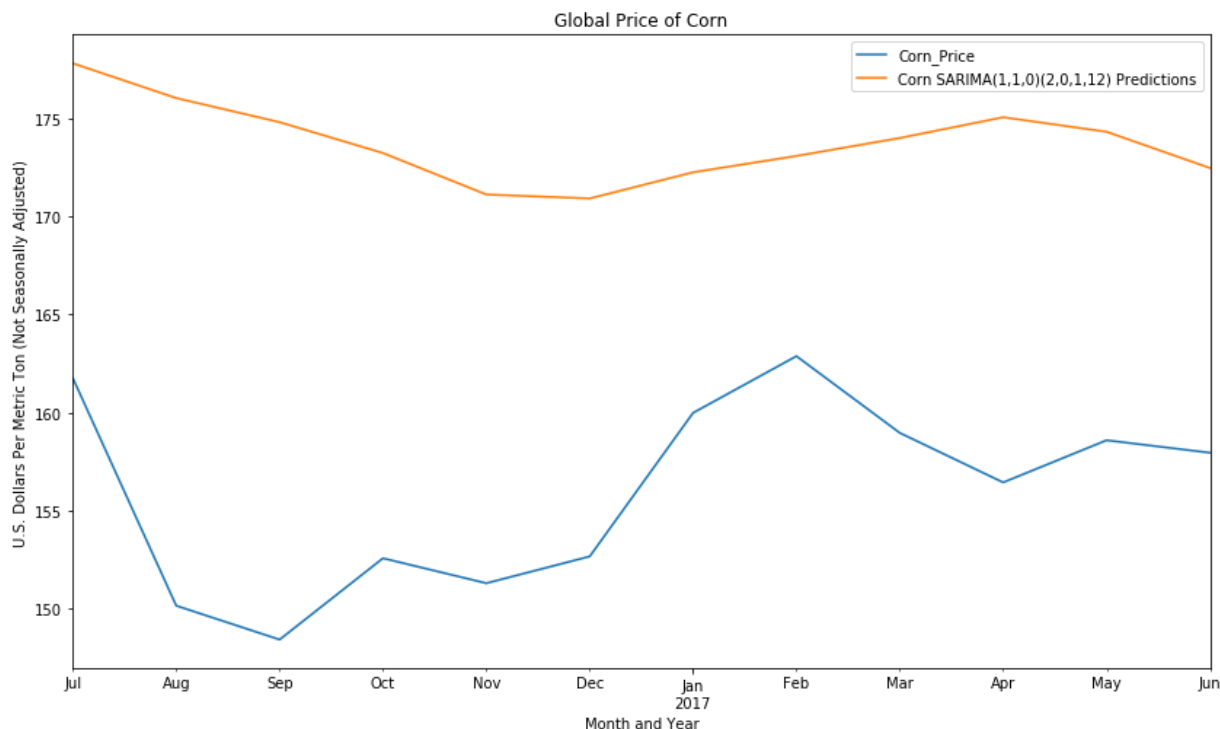
Obtain predicted values from global corn price training SARIMA(1,1,0)(2,0,1,12) model.

```
In [23]: start = len(train)
end = len(train) + len(test) - 1
corn_predictions = corn_results.predict(start=start, end=end, dynamic=False, typ:
```

Plot predictions against known global corn prices in test set.

```
In [24]: title7 = 'Global Price of Corn'
ylabel3 = 'U.S. Dollars Per Metric Ton (Not Seasonally Adjusted)'
xlabel3 = 'Month and Year'

ax3 = test['Corn_Price'].plot(legend=True, figsize=(14,8), title=title7)
corn_predictions.plot(legend=True)
ax3.autoscale(axis='x', tight=True)
ax3.set(xlabel=xlabel3, ylabel=ylabel3);
```



Evaluate global corn price training SARIMA(1,1,0)(2,0,1,12) model with mean squared error (MSE) and root mean squared error (RMSE).

```
In [25]: corn_error1 = mean_squared_error(test['Corn_Price'], corn_predictions)
corn_error2 = rmse(test['Corn_Price'], corn_predictions)
print(f'SARIMA(1,1,0)(2,0,1,12) MSE Error: {corn_error1:11.10}')
print(f'SARIMA(1,1,0)(2,0,1,12) RMSE Error: {corn_error2:11.10}')
```

SARIMA(1,1,0)(2,0,1,12) MSE Error: 338.5260711

SARIMA(1,1,0)(2,0,1,12) RMSE Error: 18.399078

Retrain SARIMA(1,1,0)(2,0,1,12) model on entire global corn price data and forecast one year into the future.

```
In [26]: corn_full_model = SARIMAX(corn_wheat2['Corn_Price'], order=(1,1,0), seasonal_order=(1,1,0,12))
corn_full_results = corn_full_model.fit()
corn_forecast = corn_full_results.predict(len(corn_wheat2), len(corn_wheat2) + 12)
```

Plot forecasted global corn prices alongside original corn price data.

```
In [27]: title8 = 'Global Price of Corn'
ylabel4 = 'U.S. Dollars Per Metric Ton (Not Seasonally Adjusted)'
xlabel4 = 'Year'

ax4 = corn_wheat2['Corn_Price'].plot(legend=True, figsize=(14,8), title=title8)
corn_forecast.plot(legend=True)
ax4.autoscale(axis='x', tight=True)
ax4.set(xlabel=xlabel4, ylabel=ylabel4);
```

